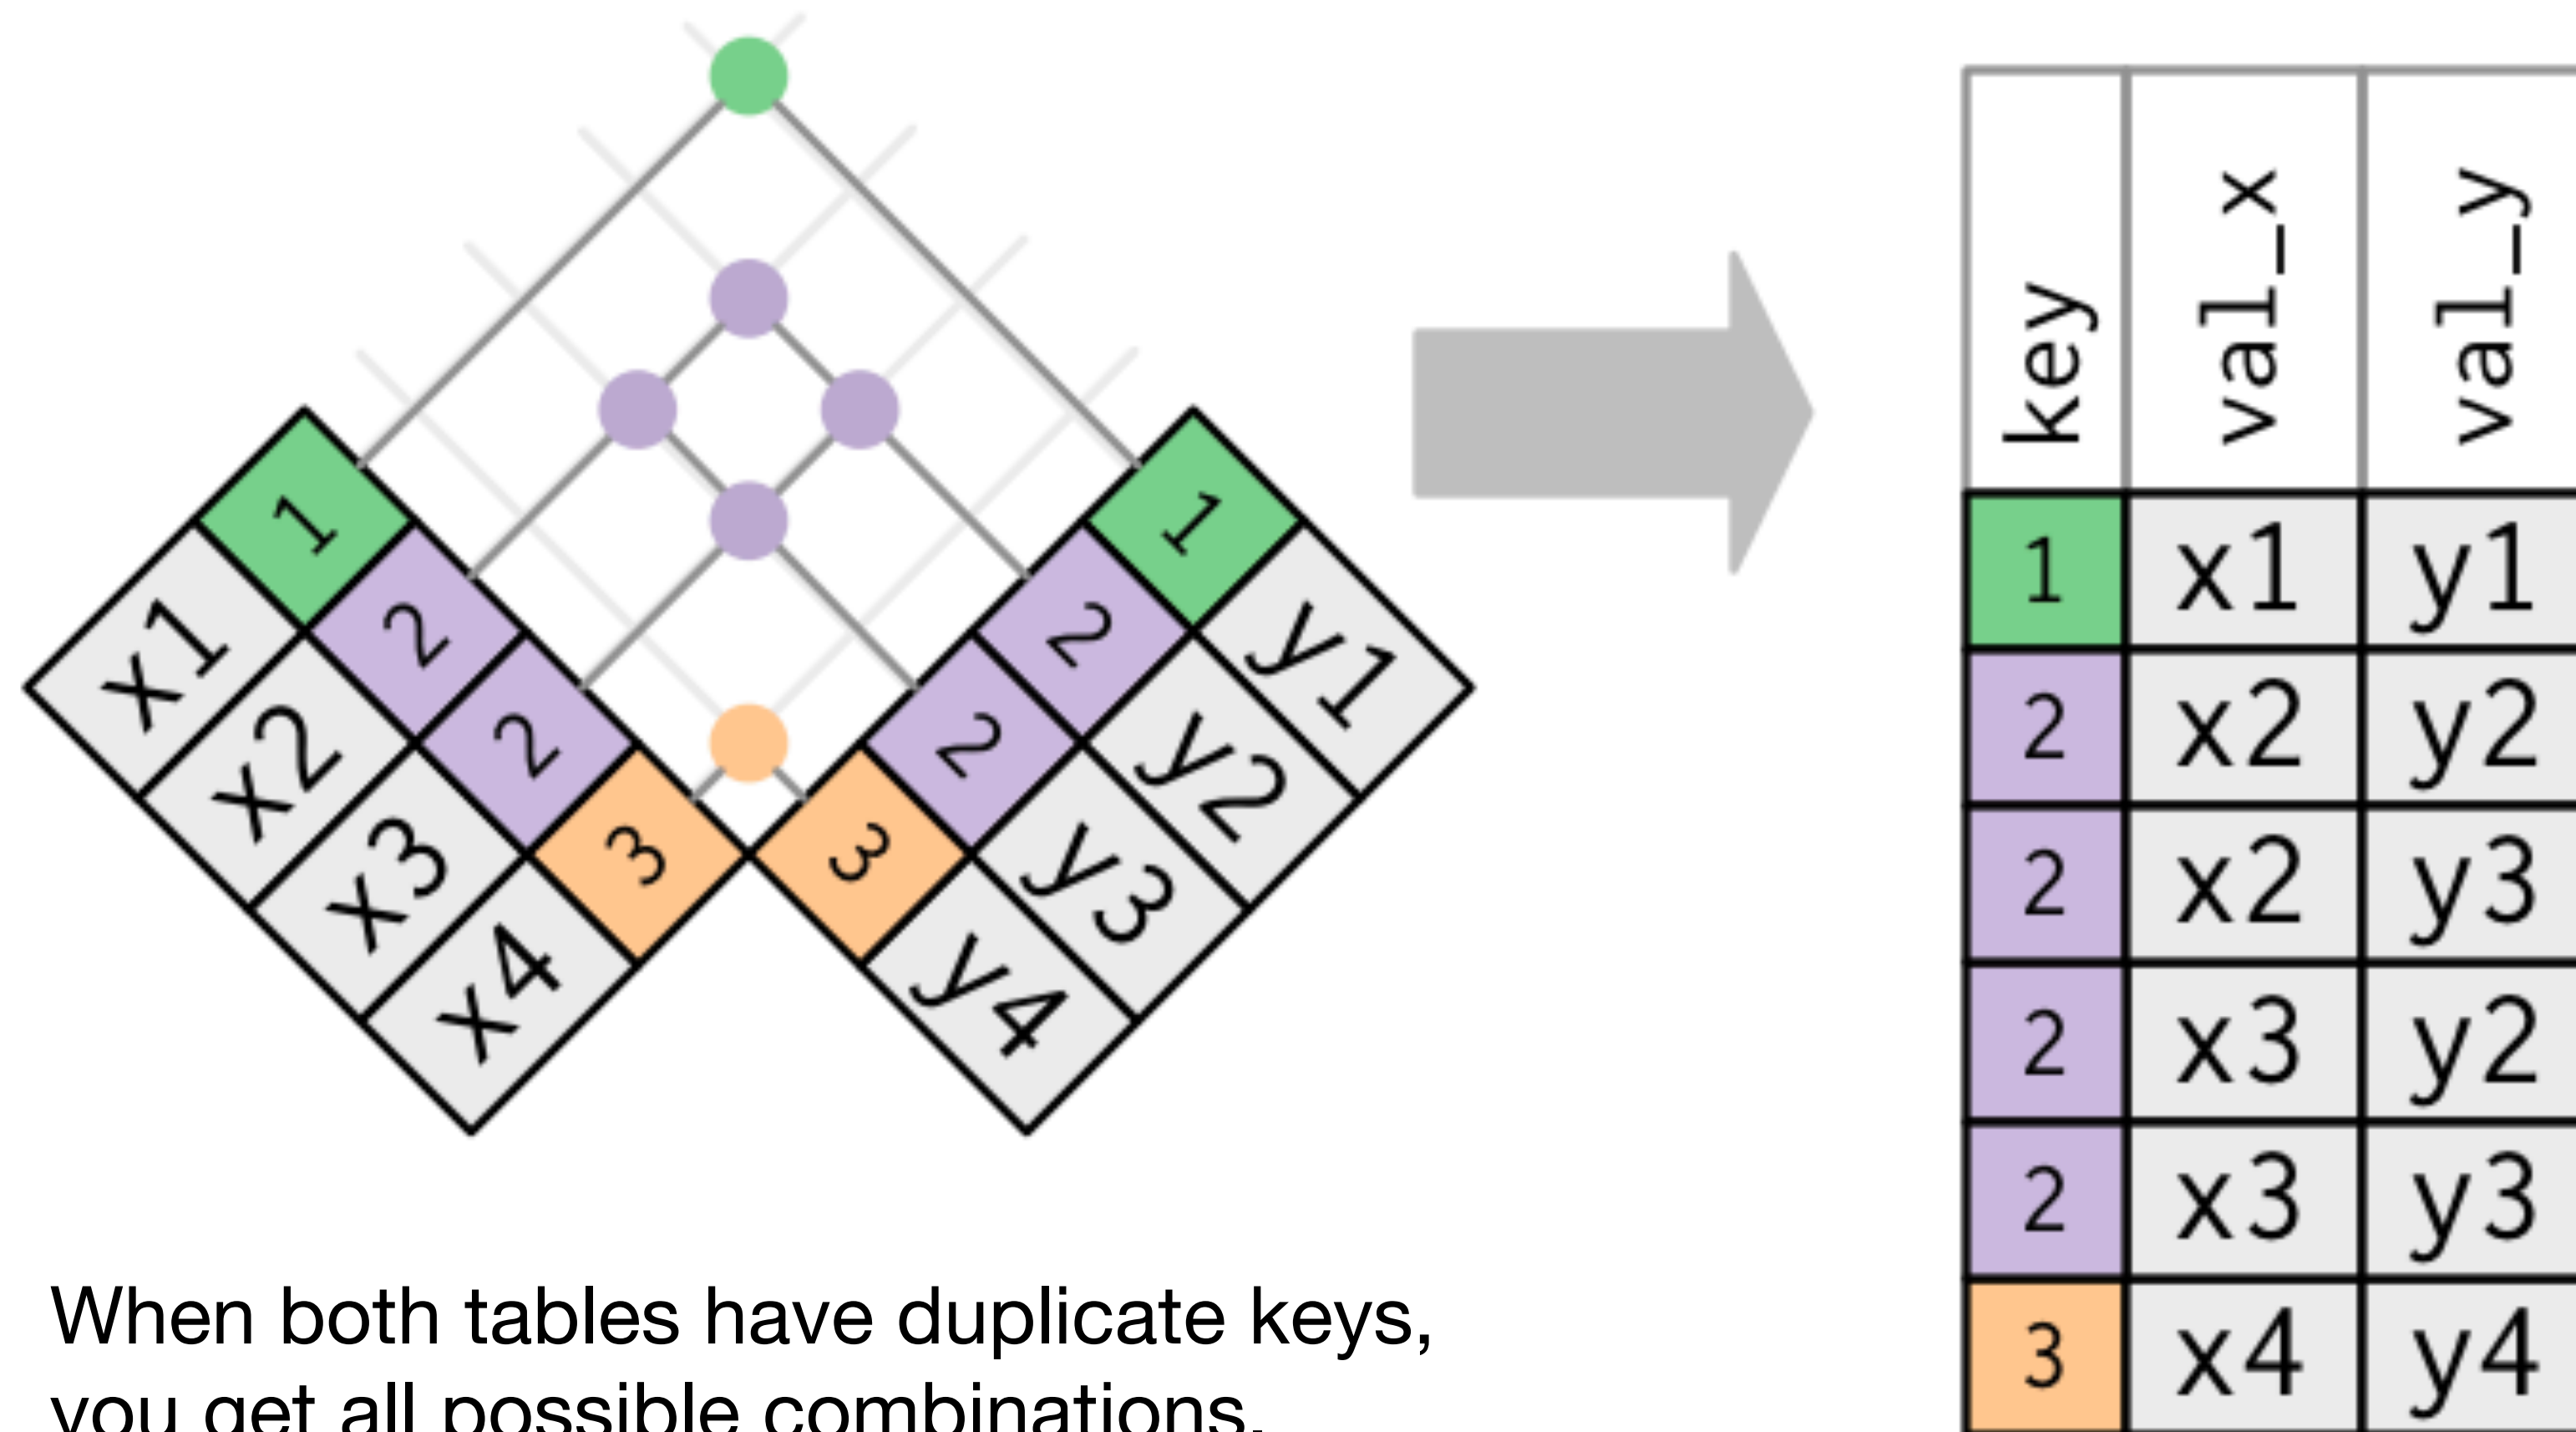


Dealing with duplicate keys can be tricky



When both tables have duplicate keys, you get all possible combinations.

Specifying keys for the match

```
flights2 %>%  
  left_join(weather)  
#> Joining, by = c("year", "month", "day", "hour", "origin")  
#> # A tibble: 336,776 x 18  
#>   year month   day hour origin dest tailnum carrier temp dewp humid  
#>   <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>    <dbl> <dbl> <dbl>  
#> 1  2013     1     1     5 EWR    IAH  N14228  UA        39.0  28.0  64.4  
#> 2  2013     1     1     5 LGA    IAH  N24211  UA        39.9  25.0  54.8  
#> 3  2013     1     1     5 JFK    MIA  N619AA  AA        39.0  27.0  61.6  
#> 4  2013     1     1     5 JFK    BQN  N804JB  B6        39.0  27.0  61.6  
#> 5  2013     1     1     6 LGA    ATL  N668DN  DL        39.9  25.0  54.8  
#> 6  2013     1     1     5 EWR    ORD  N39463  UA        39.0  28.0  64.4  
#> # ... with 3.368e+05 more rows, and 7 more variables: wind_dir <dbl>,  
#> #   wind_speed <dbl>, wind_gust <dbl>, precip <dbl>, pressure <dbl>,  
#> #   visib <dbl>, time_hour <dtm>
```

By default, dplyr functions perform a “**natural join**”, which matches observations using all variables that exist in **both** tables.