



Problem: Integrate data from multiple tables

Solution: Mutating joins

```
flights2 <- flights %>%
  select(year:day, hour, origin, dest, tailnum, carrier)
flights2
#> # A tibble: 336,776 x 8
#>   year month   day hour origin dest  tailnum carrier
#>   <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>
#> 1  2013     1     1     5 EWR    IAH   N14228   UA
#> 2  2013     1     1     5 LGA    IAH   N24211   UA
#> 3  2013     1     1     5 JFK    MIA   N619AA   AA
#> 4  2013     1     1     5 JFK    BQN   N804JB   B6
#> 5  2013     1     1     6 LGA    ATL   N668DN   DL
#> 6  2013     1     1     5 EWR    ORD   N39463   UA
#> # ... with 3.368e+05 more rows
```

```
flights2 %>%
  select(-origin, -dest) %>%
  left_join(airlines, by = "carrier")
#> # A tibble: 336,776 x 7
#>   year month   day hour tailnum carrier name
#>   <int> <int> <int> <dbl> <chr>   <chr>   <chr>
#> 1  2013     1     1     5 N14228   UA      United Air Lines Inc.
#> 2  2013     1     1     5 N24211   UA      United Air Lines Inc.
#> 3  2013     1     1     5 N619AA   AA      American Airlines Inc.
#> 4  2013     1     1     5 N804JB   B6      JetBlue Airways
#> 5  2013     1     1     6 N668DN   DL      Delta Air Lines Inc.
#> 6  2013     1     1     5 N39463   UA      United Air Lines Inc.
#> # ... with 3.368e+05 more rows
```

`left_join()` matches two tables based on their keys and adds the new variables from the second one to the first one

# Problem: Integrate data from multiple tables

## Solution: Mutating joins

```
flights2 <- flights %>%
  select(year:day, hour, origin, dest, tailnum, carrier)
flights2
#> # A tibble: 336,776 x 8
#>   year month   day hour origin dest tailnum carrier
#>   <int> <int> <int> <dbl> <chr>  <chr> <chr>   <chr>
#> 1  2013     1     1     5 EWR    IAH  N14228  UA
#> 2  2013     1     1     5 LGA    IAH  N24211  UA
#> 3  2013     1     1     5 JFK    MIA  N619AA  AA
#> 4  2013     1     1     5 JFK    BQN  N804JB  B6
#> 5  2013     1     1     6 LGA    ATL  N668DN  DL
#> 6  2013     1     1     5 EWR    ORD  N39463  UA
#> # ... with 3.368e+05 more rows
```

```
flights2 %>%
  select(-origin, -dest) %>%
  left_join(airlines, by = "carrier")
#> # A tibble: 336,776 x 7
#>   year month   day hour tailnum carrier name
#>   <int> <int> <int> <dbl> <chr>   <chr>  <chr>
#> 1  2013     1     1     5 N14228  UA      United Air Lines Inc.
#> 2  2013     1     1     5 N24211  UA      United Air Lines Inc.
#> 3  2013     1     1     5 N619AA  AA      American Airlines Inc.
#> 4  2013     1     1     5 N804JB  B6      JetBlue Airways
#> 5  2013     1     1     6 N668DN  DL      Delta Air Lines Inc.
#> 6  2013     1     1     5 N39463  UA      United Air Lines Inc.
#> # ... with 3.368e+05 more rows
```

`left_join()` matches two tables based on their keys and adds the new variables from the second one to the first one

# There are several types of join operations

```
x <- tribble(
  ~key, ~val_x,
  1, "x1",
  2, "x2",
  3, "x3"
)
y <- tribble(
  ~key, ~val_y,
  1, "y1",
  2, "y2",
  4, "y3"
)
```

